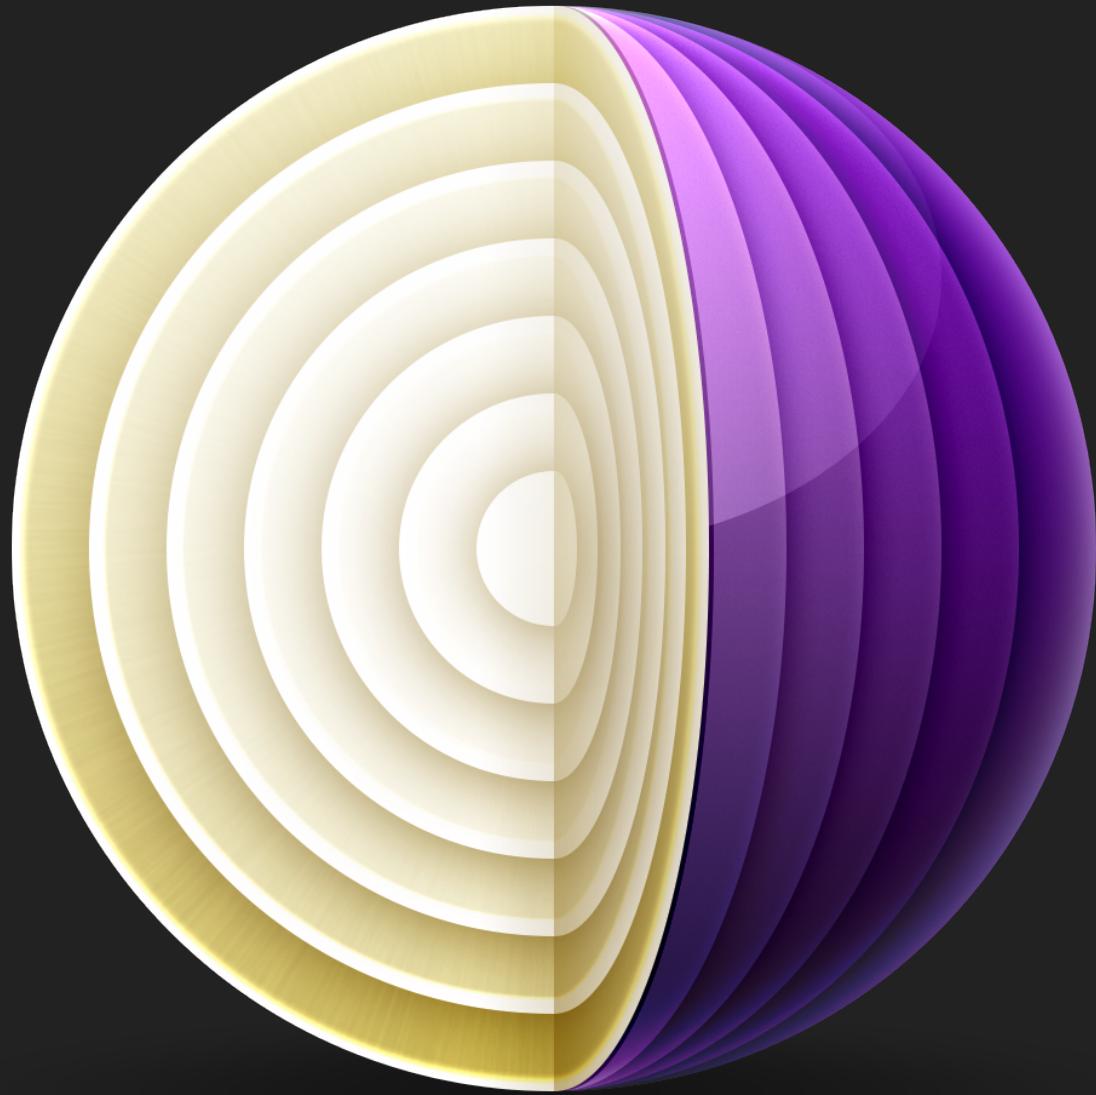


ITSECX -  @FREDERICJACOBS

REFLECTIONS ON TRUSTING
BITCODE

~ >>> WHOIS FREDERIC



iCepa



Signal



Both Open-Source Software

Talk based on blog post

The screenshot shows a web browser window with a blue header bar. The title bar reads "Why I'm not enabling Bitcode" and "A Medium Corporation [US] https://medium.com/@FredericJacobs/why-i-m-not-enabling-bitcode-f35cd8fbfcc5". The main content area is a white page from Medium. At the top left is the Medium logo. To the right is a search bar with "Search Medium" and a "Sign in / Sign up" button. Below the search bar, there's a profile picture of Frederic Jacobs and his name. A timestamp "Sep 24 · 14 min read" is also present. The main title "Why I'm not enabling Bitcode" is in large bold black font. Below it, a subtitle "Thoughts on application binaries packaging and software distribution" is in smaller gray font. The main text begins with "At Apple's annual WWDC developer event, the compiler infrastructure team unveiled ‘Bitcode’ and recommended iOS developers to opt-in , even requiring it for Apple Watch applications." A section titled "The Problem with fat archives" follows, with text explaining the nature of fat archives and how they are handled by the system.

Why I'm not enabling Bitcode

Thoughts on application binaries packaging and software distribution

At Apple's annual WWDC developer event, the compiler infrastructure team unveiled “**Bitcode**” and recommended iOS developers to opt-in , even requiring it for Apple Watch applications.

The Problem with fat archives

Multi-architecture iOS (and OS X) executables are packaged in fat archives. The fat format (not to same as the filesystem) is quite simple: the application binaries contain headers describing the scope of the instructions for a given infrastructure. Fat headers of a binary can be looked up by running:

[HTTPS://MEDIUM.COM/@FREDERICJACOBS/WHY-I-M-NOT-ENABLING-BITCODE-F35CD8FBFCC5](https://medium.com/@FredericJacobs/why-i-m-not-enabling-bitcode-f35cd8fbfcc5)

MOTIVATION

Understanding the impact of
Bitcode on the distribution of
iOS security-sensitive software.

OR HOW MORGAN PUT IT

Morgan Marquis-Boire [Following](#)

@headhntr

Excellent, sprawling post on Apple's BitCode announcement, XCodeGhost, and supply chain security medium.com/@FredericJacob ... by [@FredericJacobs](https://medium.com/@FredericJacobs)

```
fat_magic 0xc0fe0000  
nfat_arch 2  
architecture 0  
    cputype 12  
    cpusubtype 9  
    capabilities 0x0  
    offset 16384  
    size 5306144  
    align 2^14 (16384)  
architecture 1  
    cputype 16777228  
    cpusubtype 0  
    capabilities 0x0  
    offset 5324800  
    size 6457600  
    align 2^14 (16384)
```

[>>>](https://medium.com/@FredericJacobs/~/M/i/M/S/P/Signal.app)

Why I'm not enabling Bitcode
Thoughts on application binaries packaging and software distribution
[medium.com](https://medium.com/@FredericJacobs)

OUTLINE



Software Distribution



Bitcode



Security Implications

SOFTWARE DISTRIBUTION



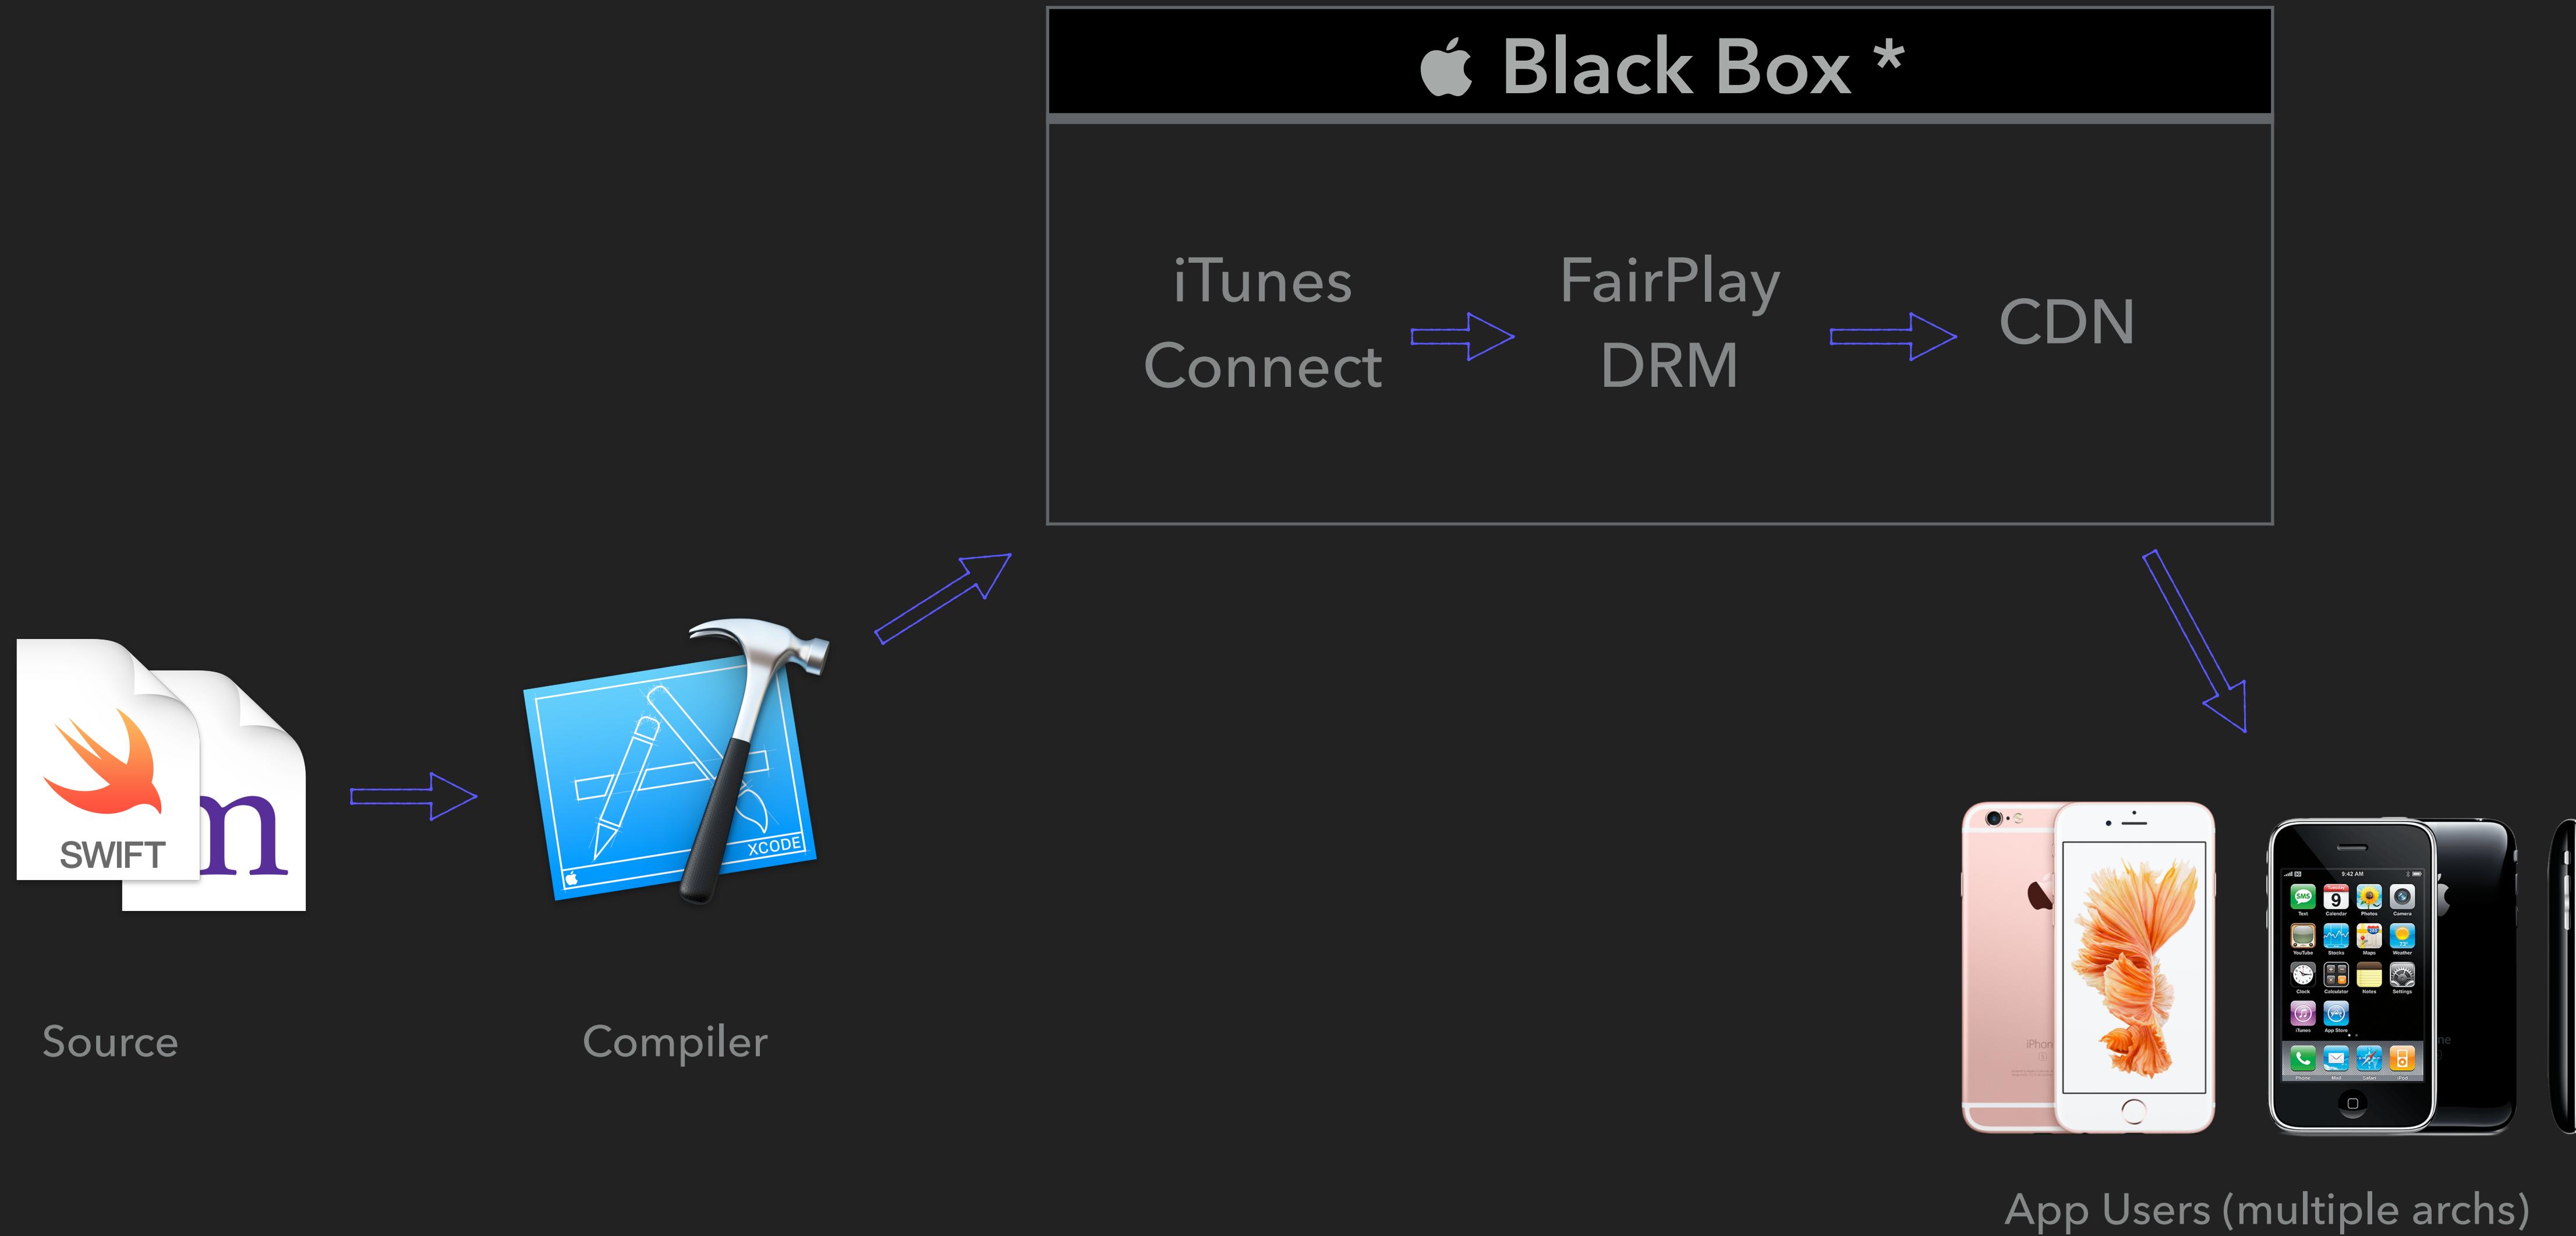
The EFI monster! [@osxreverser](https://twitter.com/osxreverser)

The App Stores promised you a security gate. The reality is that you get shit and a big black box you can't control and audit ;-)

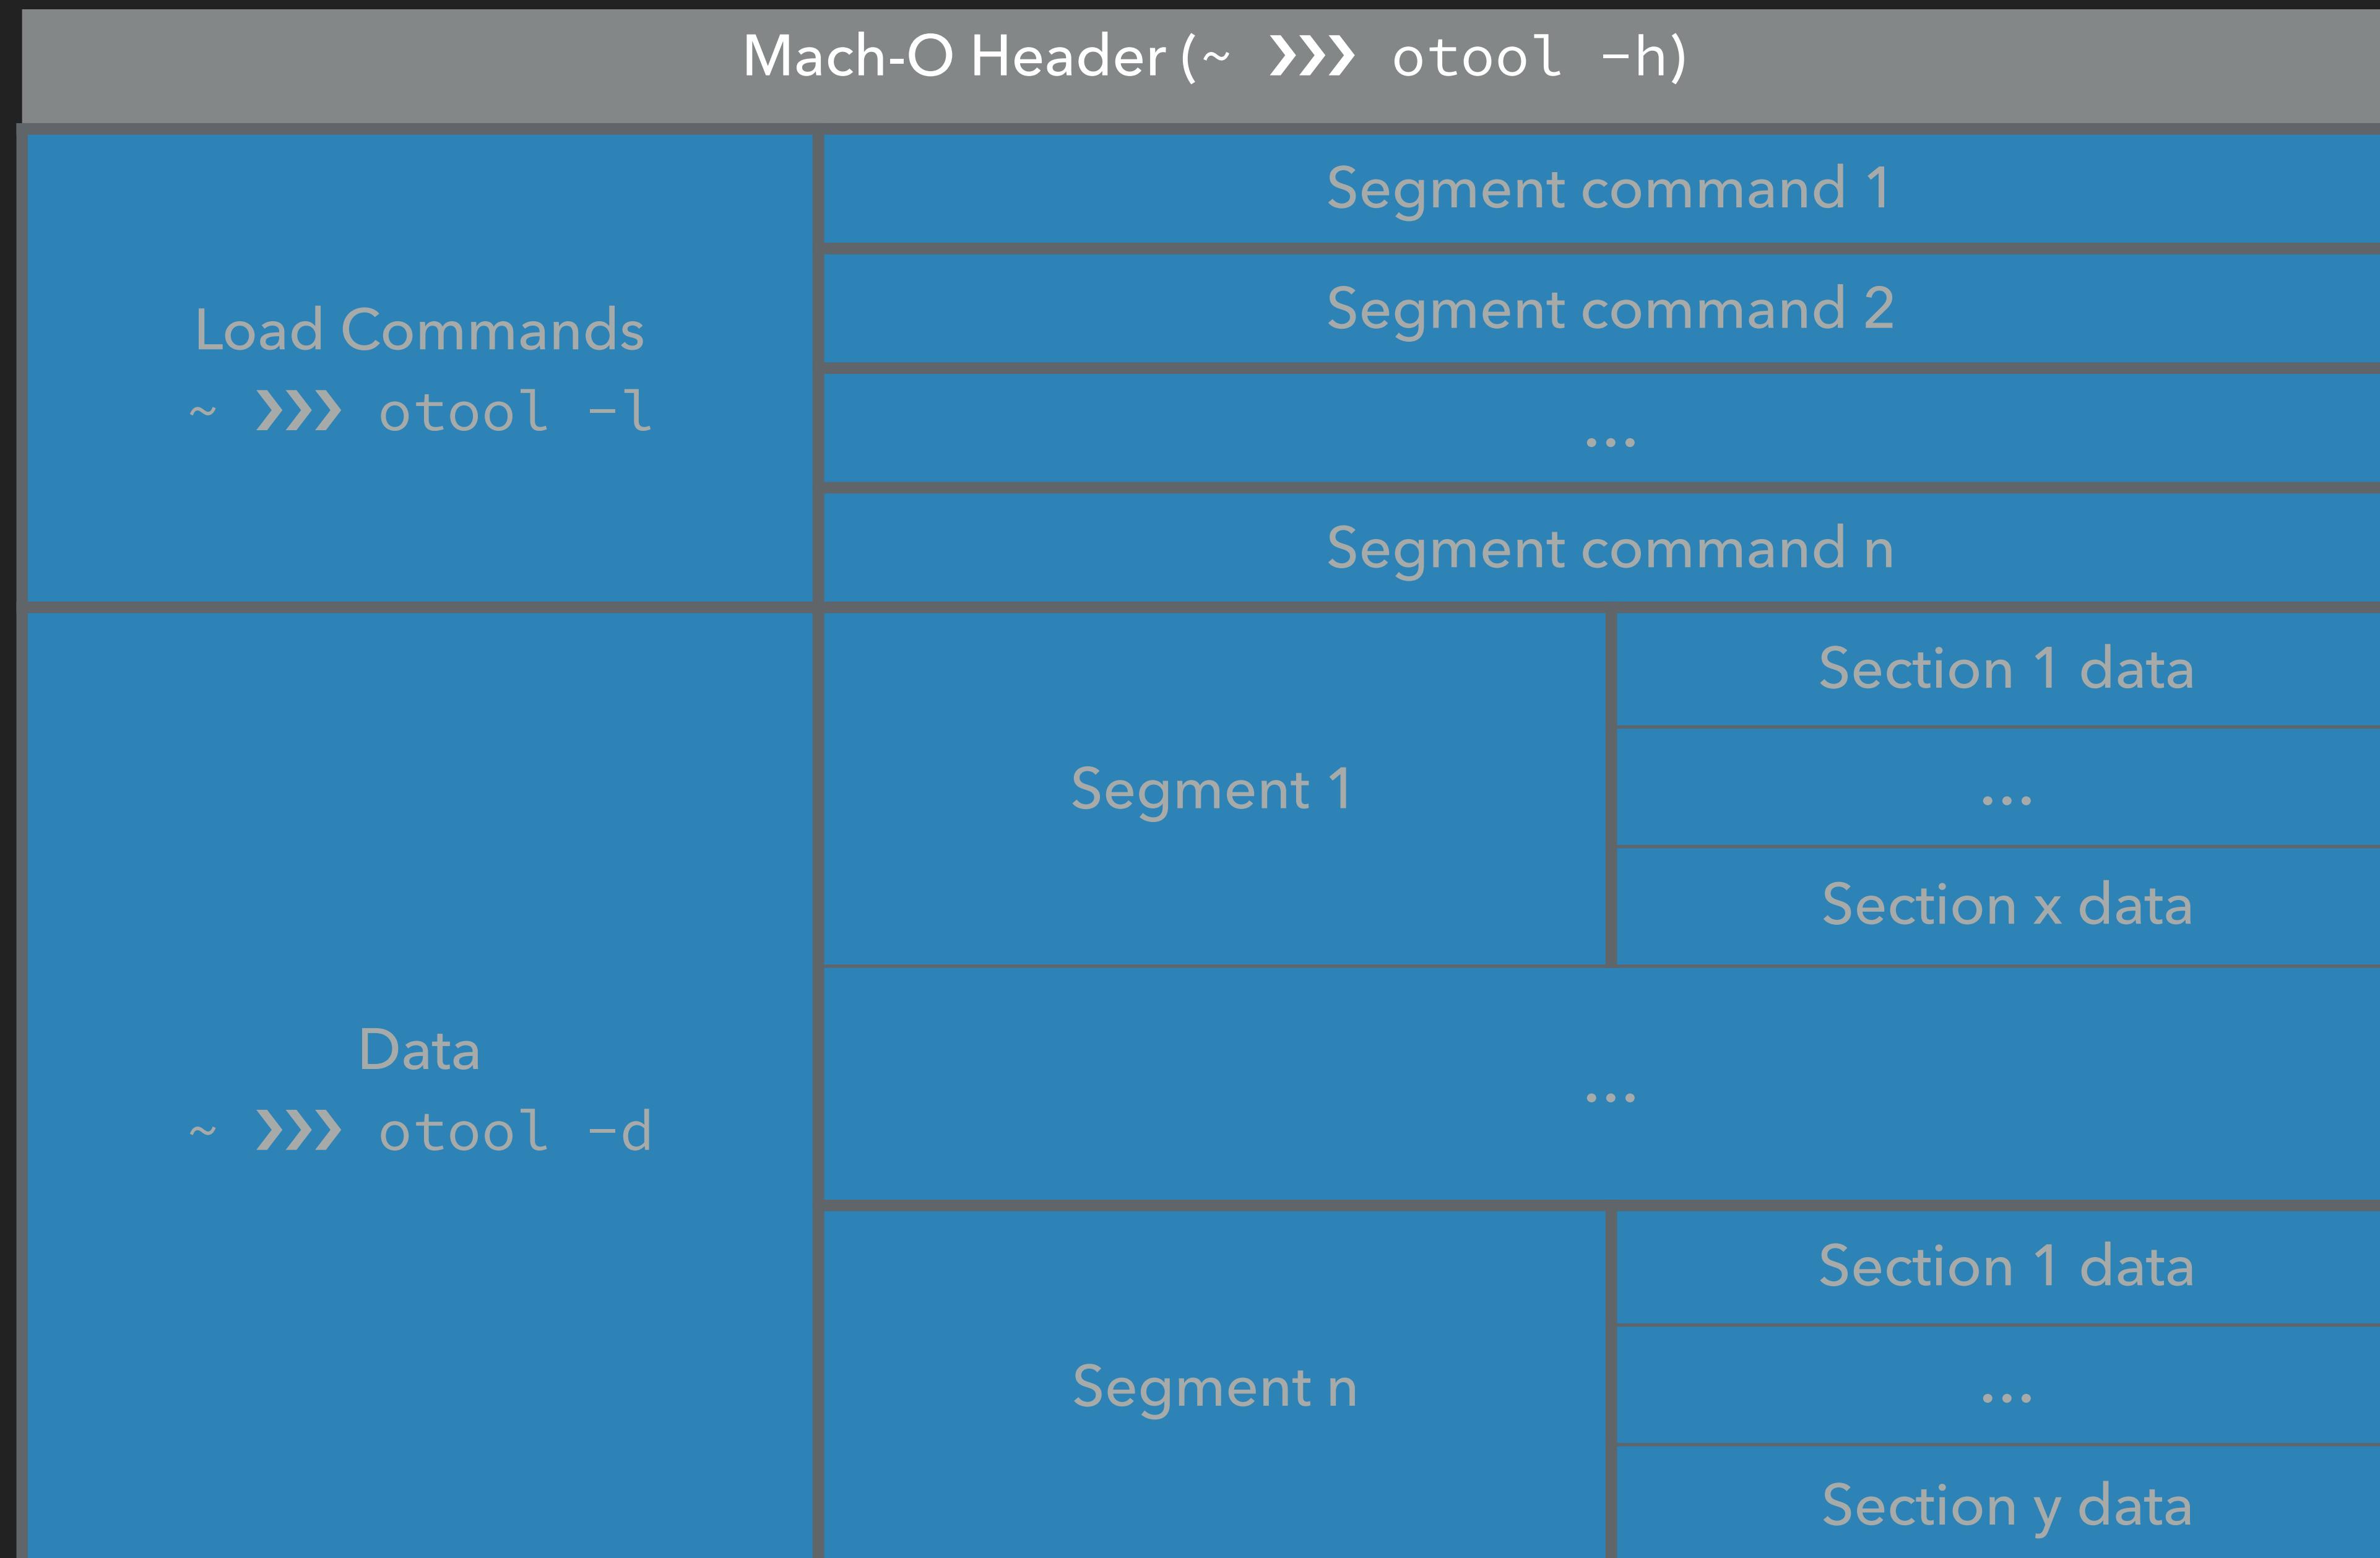
RETWEETS 2 LIKES 4

9:43 AM - 4 Nov 2015

© 2015 Twitter About Help Terms Privacy Cookies Ads info



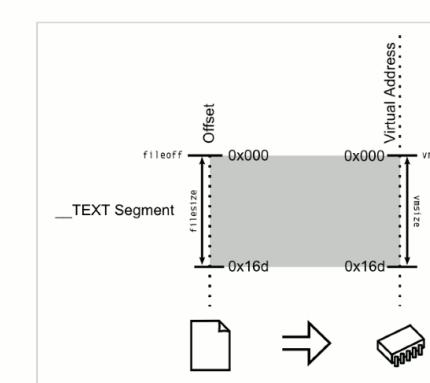
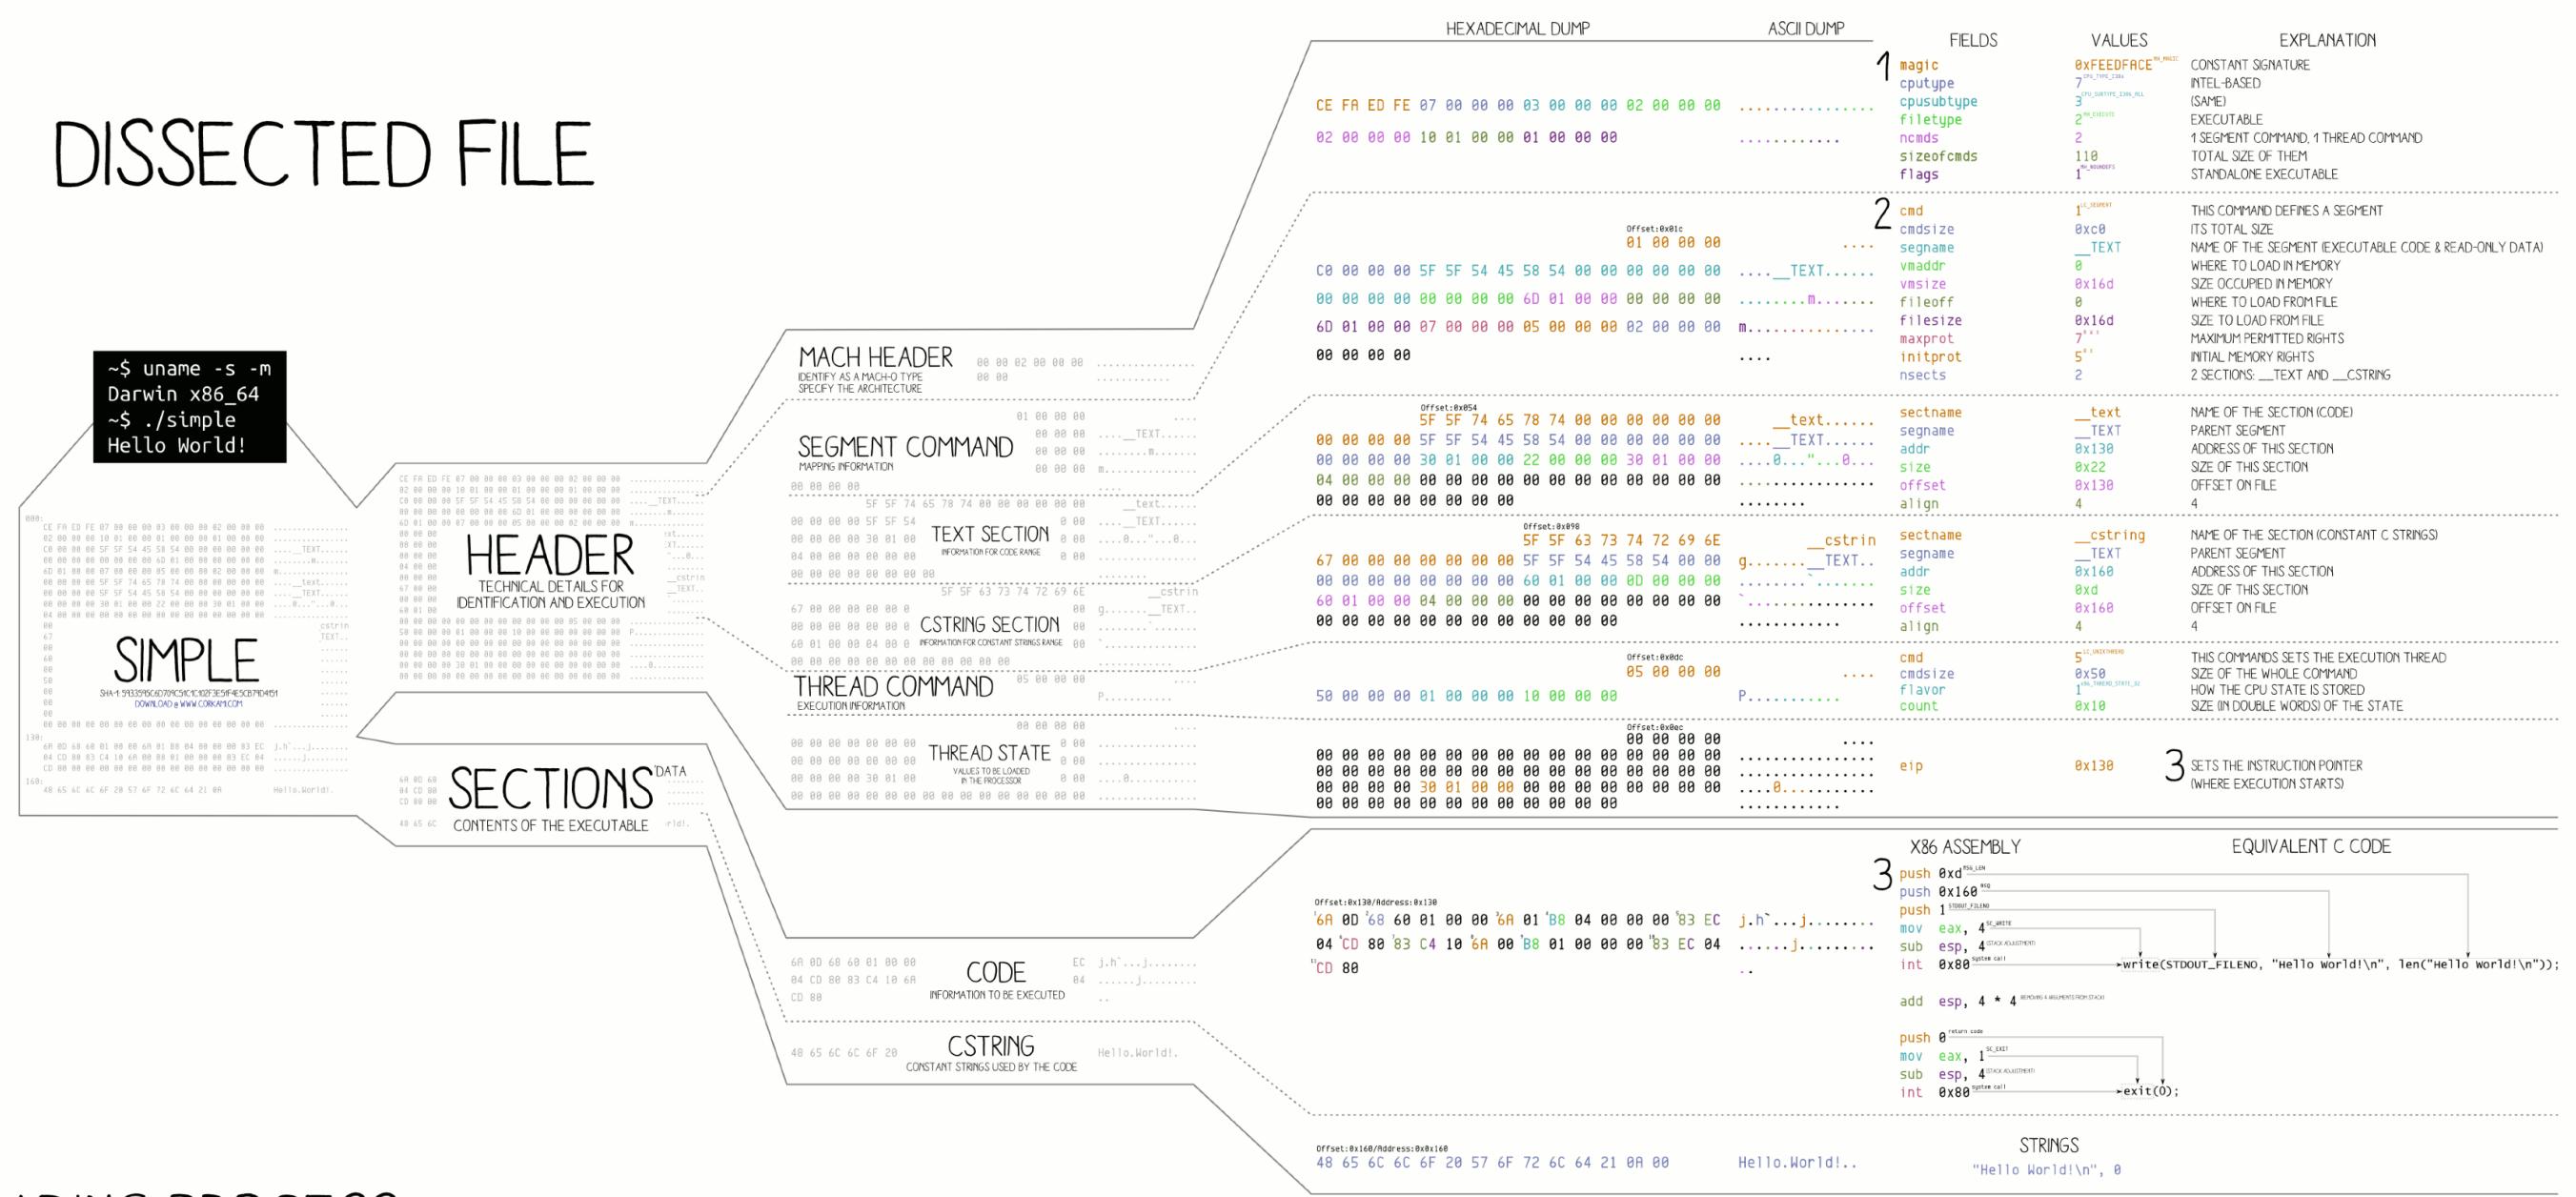
* Heh, it's a black box, I actually have no clue what's really in there but given functionality these components must be there.



Detailed Mach-O Description

MACH-O¹⁰¹ an OS X executable walkthrough ANGE ALBERTINI CORKAMI.COM

DISSECTED FILE

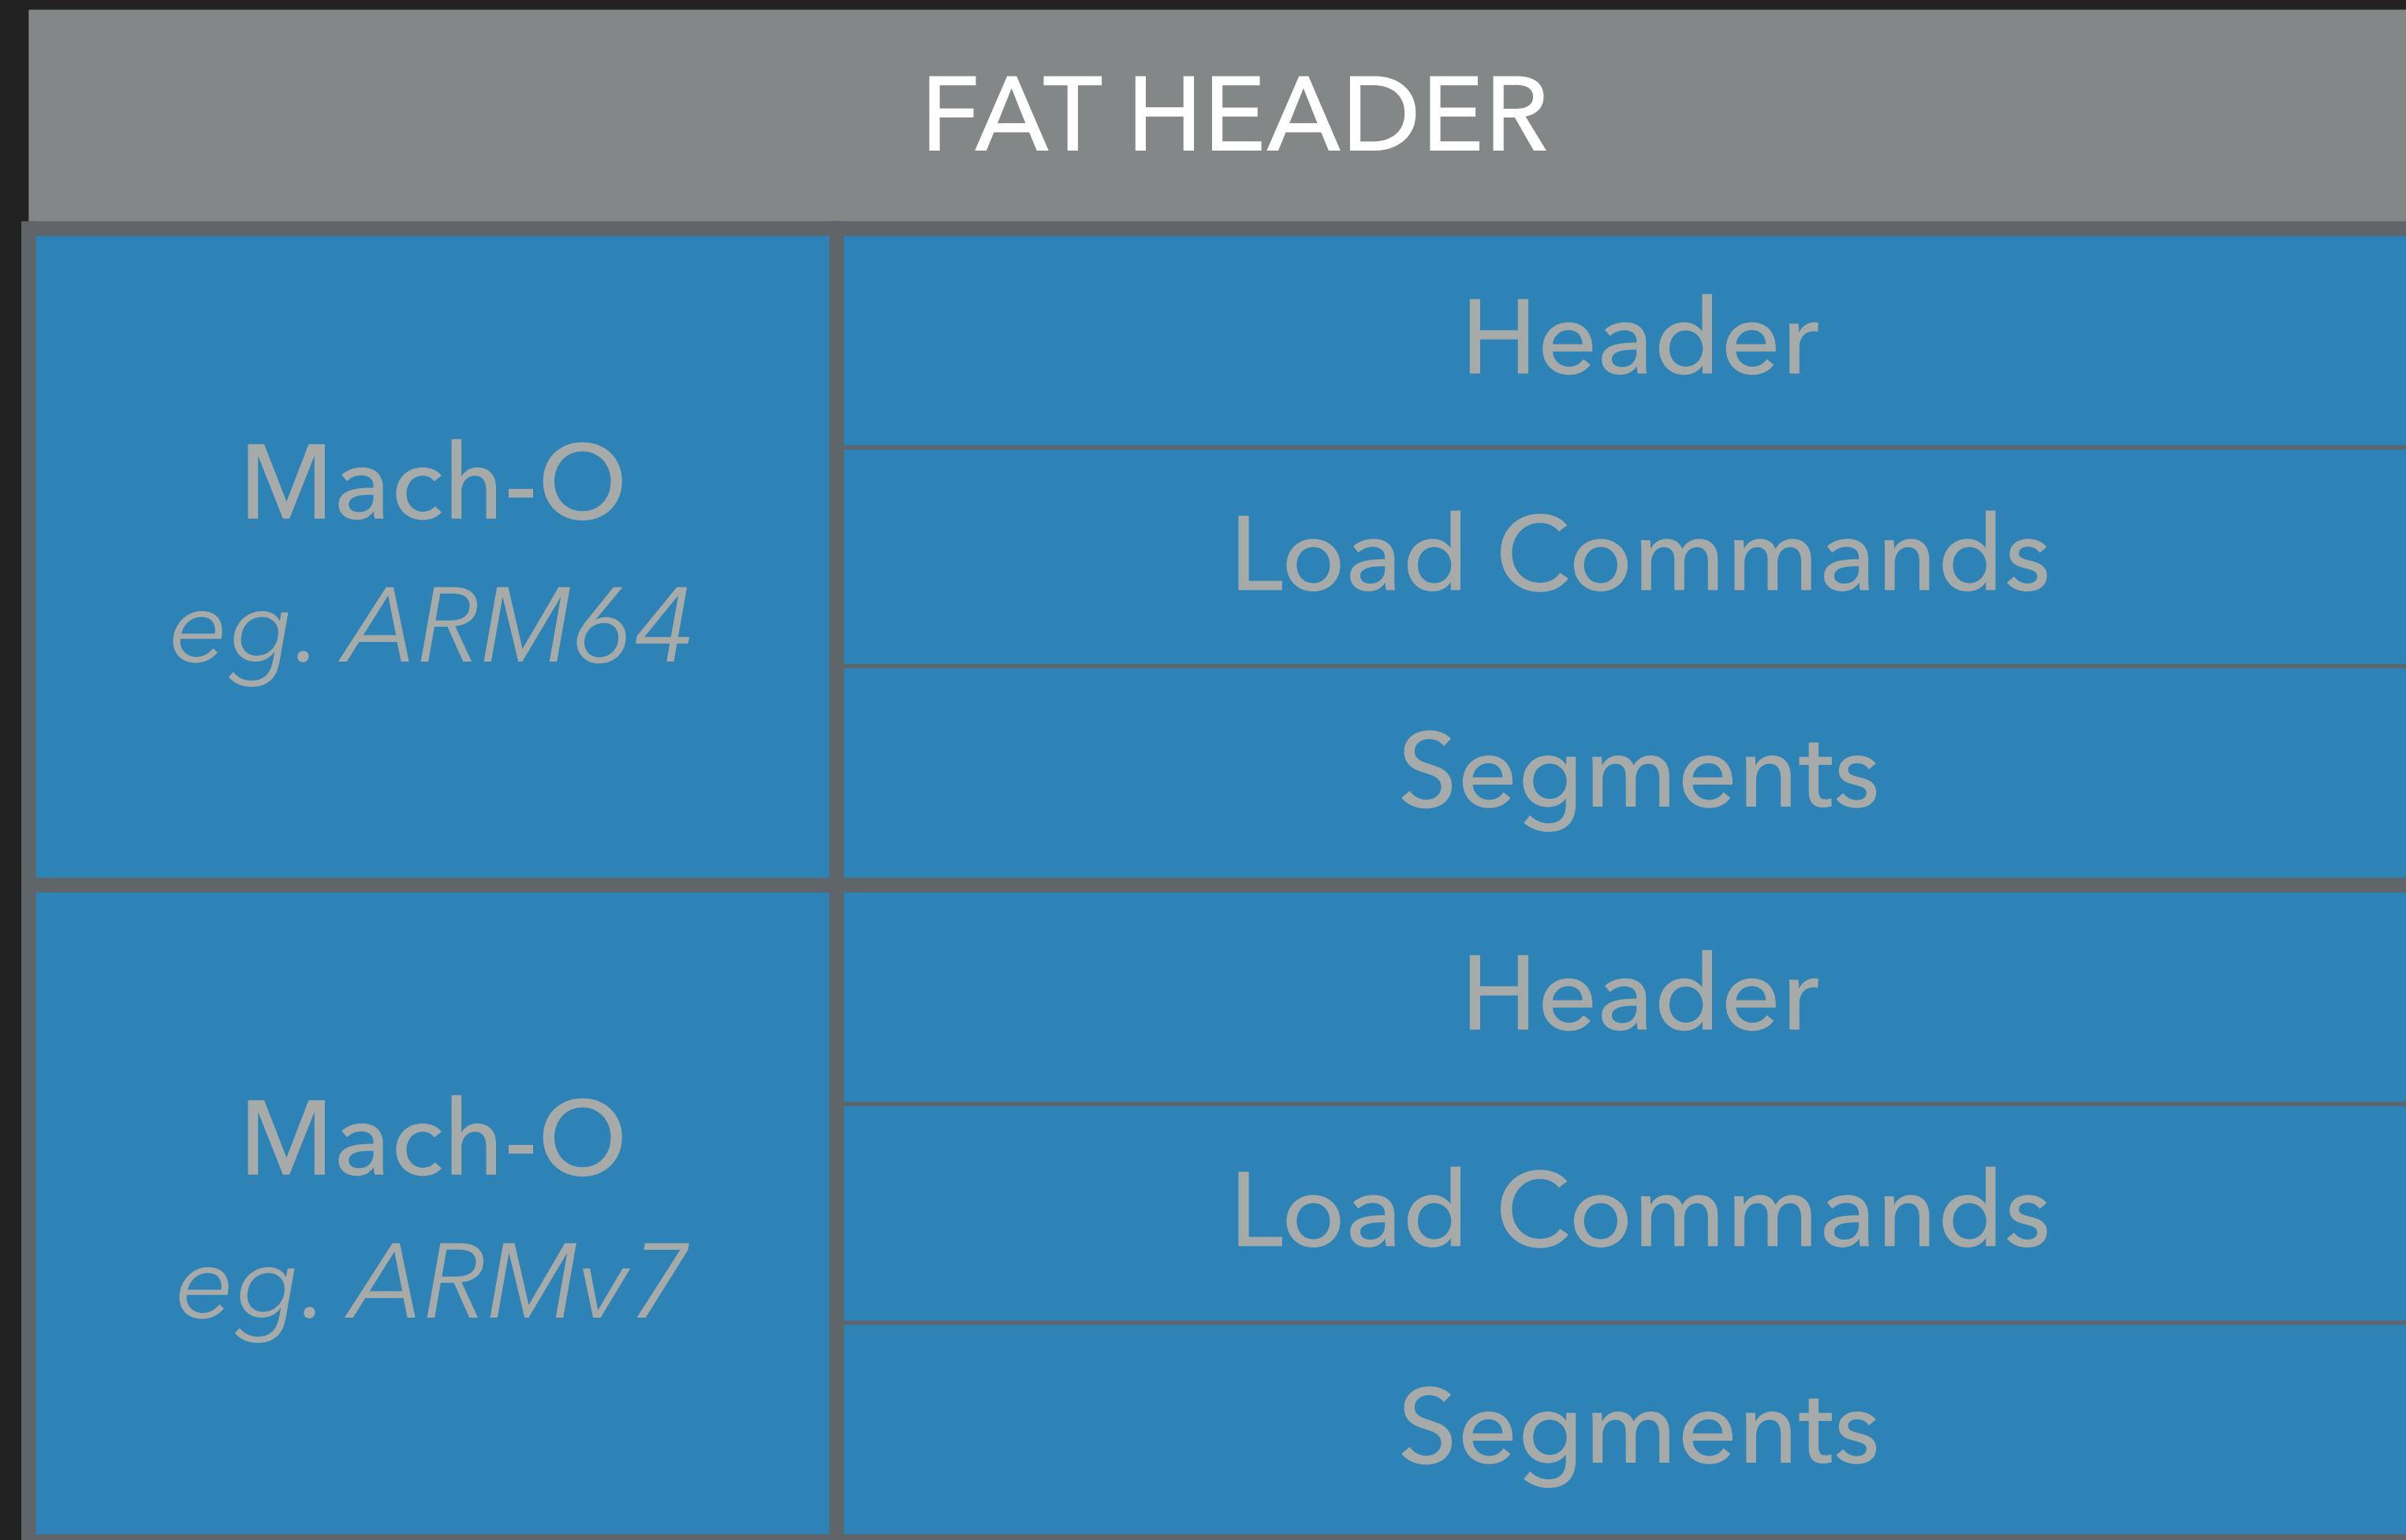


TRIVIA

THE MACH-O FORMAT COMES FROM THE MACH KERNEL, CREATED AT CMU, IN 1985

IT'S USED AMONG OTHERS,
BY OS X, IOS, STEP...
ON IPHONES, IPODS, MACS...

UNIVERSAL MACH-O FILES



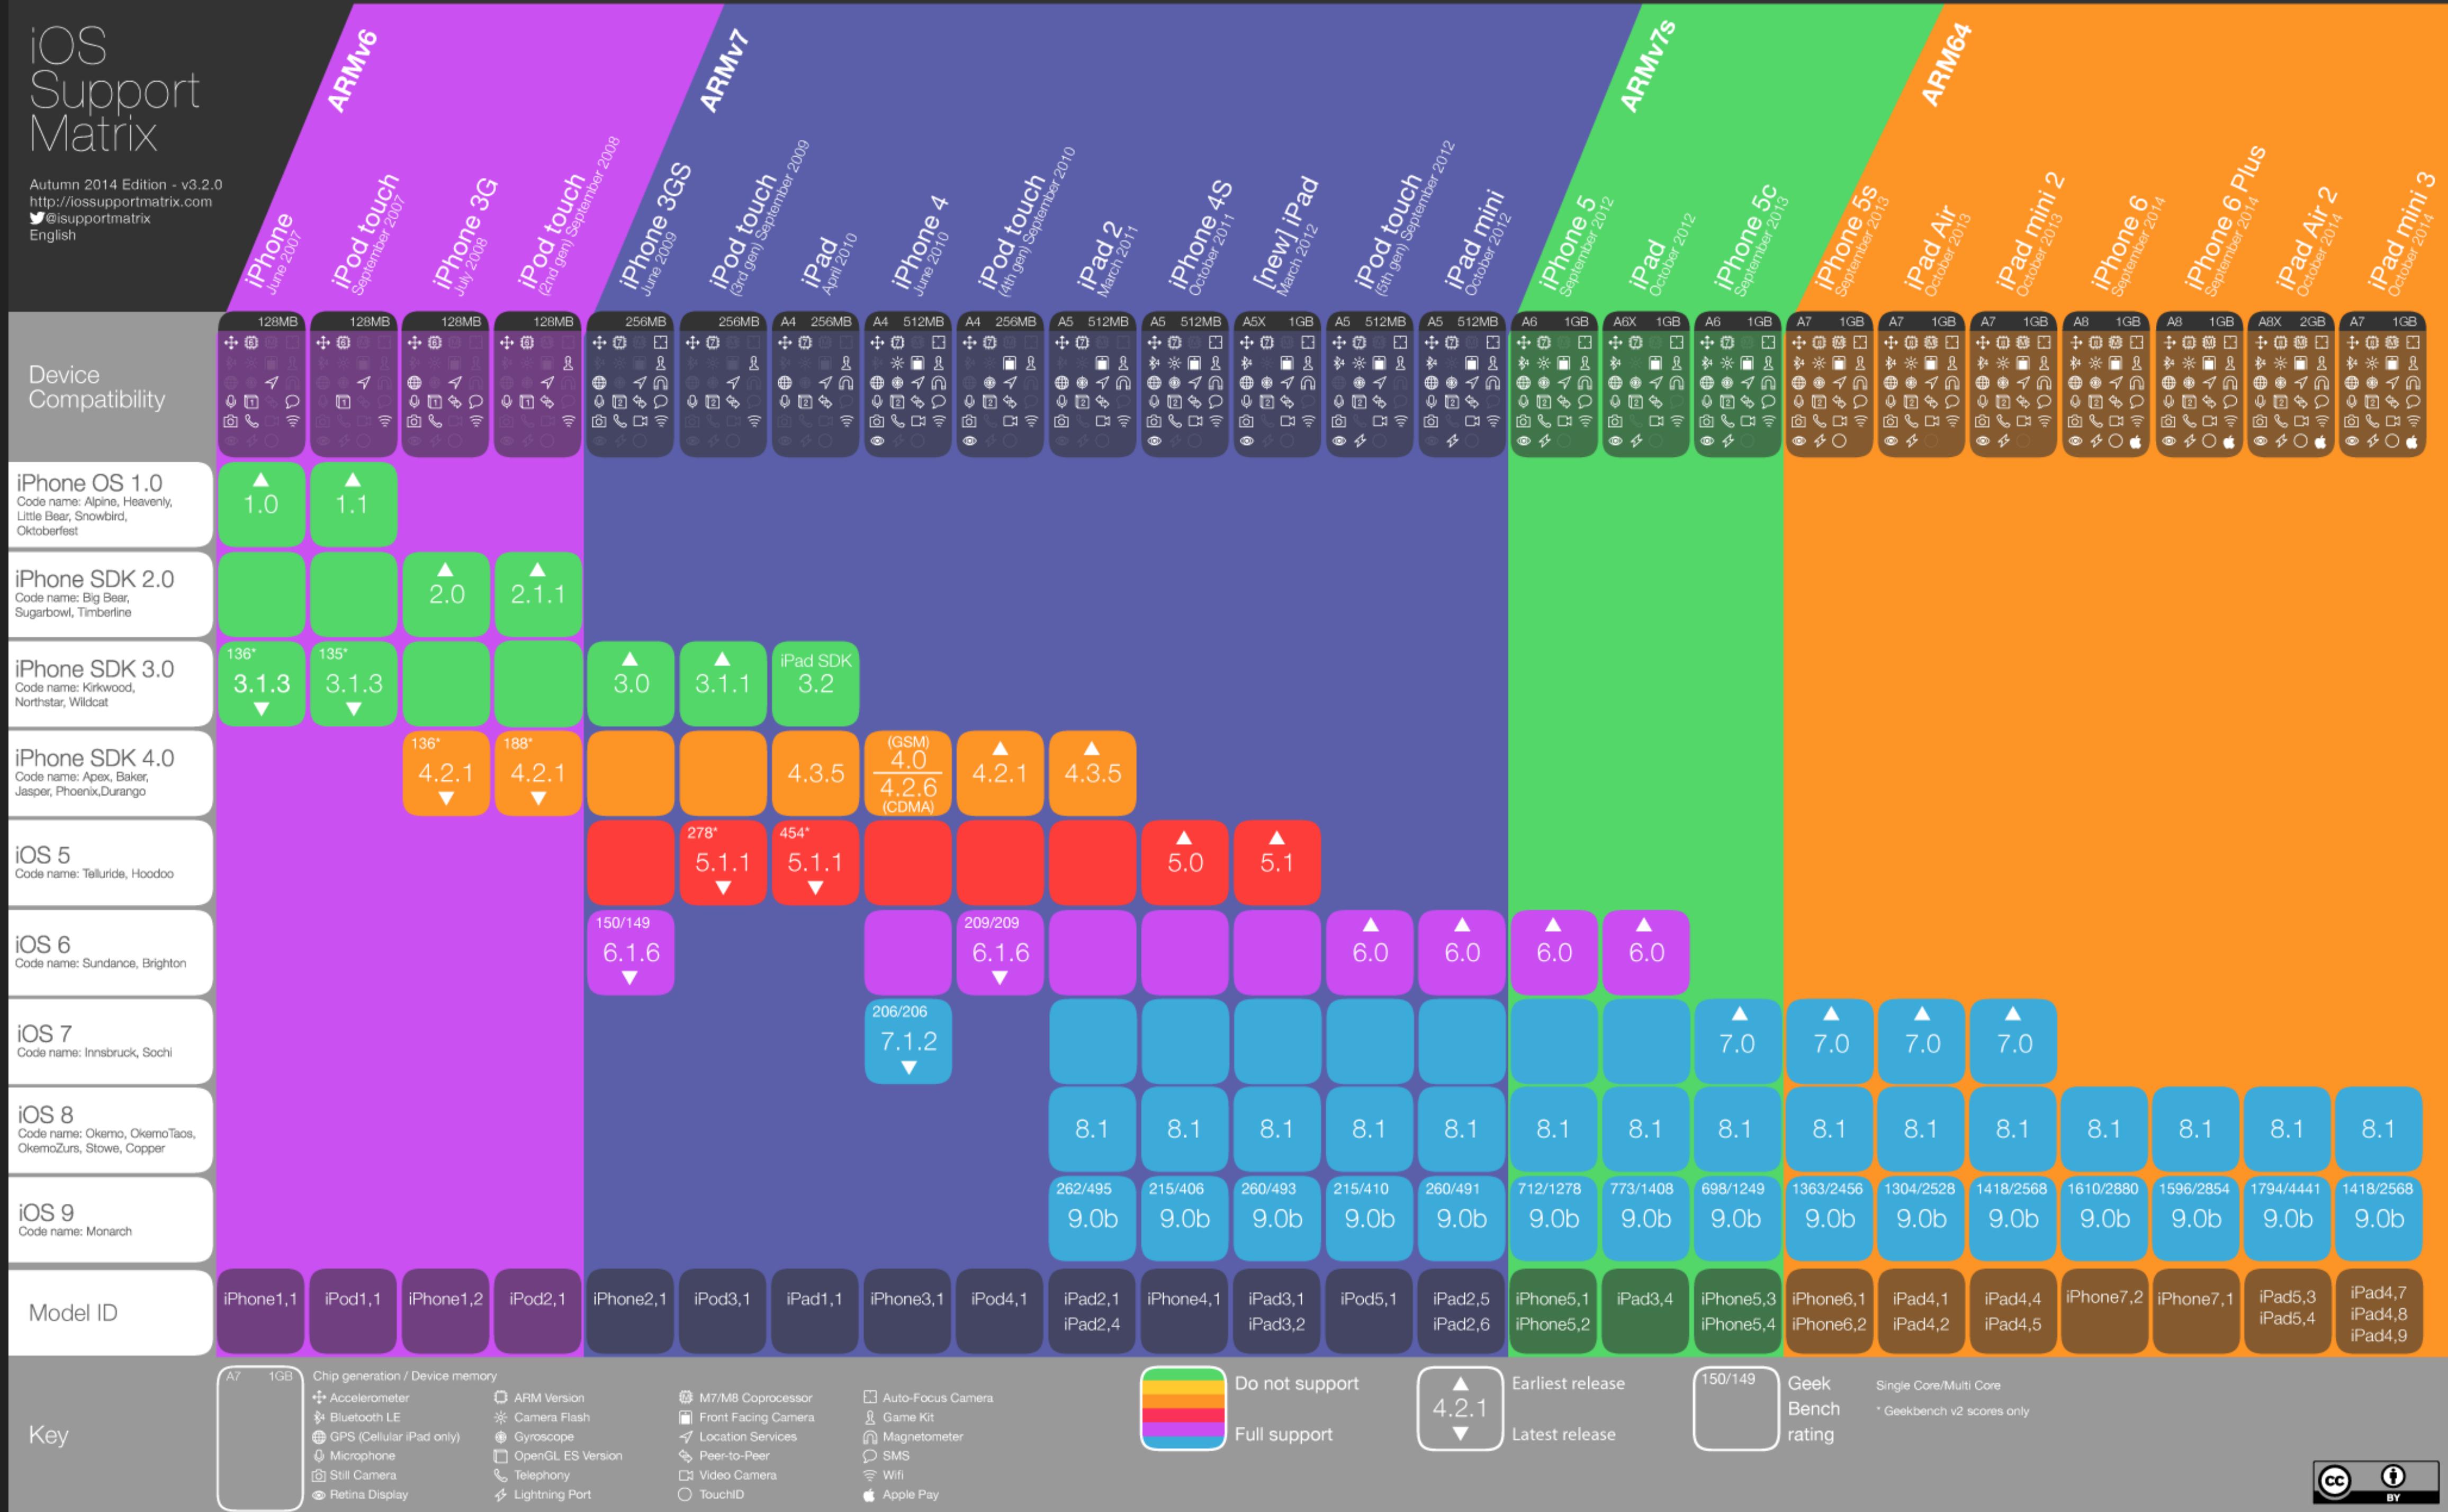
FAT HEADERS LOOKUP

```
1. zsh
~/M/i/i/M/S/Signal.app >>> otool -f Signal_Binary_iTunes
Fat headers
fat_magic 0xcafebabe
nfat_arch 2
architecture 0
    cputype 12
    cpusubtype 9
    capabilities 0x0
    offset 16384
    size 5306144
    align 2^14 (16384)
architecture 1
    cputype 16777228
    cpusubtype 0
    capabilities 0x0
    offset 5324800
    size 6457600
    align 2^14 (16384)
~/M/i/i/M/S/Signal.app >>>
```

- nfat_arch: Number of supported architectures
- fat_magic: Defines configuration: 32 vs 64 bits, endianness
- armv7: ARM has CPU type 12, armv7 is CPU subtype 9. (The armv7-specific instructions start at offset 16384 and go on for the next 5306144 bytes.)
- arm64: CPU type 16777228, the ARM 64-bit CPU type.

iOS Support Matrix

Autumn 2014 Edition - v3.2.0
<http://iossupportmatrix.com>
@isupportmatrix
English



UNIVERSAL BINARIES

- ▶ Advantages
 - ▶ Supports multiple (micro-)architectures in a single file
 - ▶ User can back up older iPhone and restore on device with newer architecture without having to download any additional material
- ▶ Disadvantages
 - ▶ They are fat (yeah, they are literally called **fat binaries**)
Adding new architectures comes at an important storage cost

ios 9

App Thinning

- ▶ Goal: Reduce App Store downloads size
- ▶ How? Only distribute to the user necessary resources according to screen resolution and device architecture
- ▶ No recompilation



BITCODE

LLVM Frontend

```
~ >>> cat helloworld.c

/* Hello World program */

#include<stdio.h>

main()
{
    printf("Hello World");
}
```

```
~ >>> clang helloworld.c
```

LLVM-IR

Wait, why does it specify a target?

- ▶ LLVM IR isn't target independent unlike what is widely believed.
- ▶ Clang source languages (ObjC, C, C++ ...) semantics are not target independent, so IR can't be target independent.

```
; ModuleID = 'helloworld.c'
target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-
S128"
target triple = "x86_64-apple-macosx10.11.0"

@.str = private unnamed_addr constant [12 x i8] c"Hello
World\00", align 1

; Function Attrs: nounwind ssp uwtable
define i32 @main() #0 {
    %1 = call i32 (i8*, ...)* @printf(i8* getelementptr
inbounds ([12 x i8]* @.str, i32 0, i32 0))
    ret i32 0
}

declare i32 @printf(i8*, ...) #1

attributes #0 = { nounwind ssp uwtable "less-precise-
fpmad"="false" "no-frame-pointer-elim"="true" "no-
frame-pointer-elim-non-leaf" "no-infs-fp-math"="false"
"no-nans-fp-math"="false" "stack-protector-buffer-
size"="8" "target-cpu"="core2" "target-
features"="+ssse3,+cx16,+sse,+sse2,+sse3" "unsafe-fp-
math"="false" "use-soft-float"="false" }
attributes #1 = { "less-precise-fpmad"="false" "no-
frame-pointer-elim"="true" "no-frame-pointer-elim-non-
leaf" "no-infs-fp-math"="false" "no-nans-fp-
math"="false" "stack-protector-buffer-size"="8"
"target-cpu"="core2" "target-
features"="+ssse3,+cx16,+sse,+sse2,+sse3" "unsafe-fp-
math"="false" "use-soft-float"="false" }

!llvm.module.flags = !{!0}
!llvm.ident = !{!1}

!0 = !{i32 1, !"PIC Level", i32 2}
!1 = !{!"Apple LLVM version 7.0.0 (clang-700.1.76)"}
```

LLVM Backend

- ▶ Compiling LLVM to target binary
- ▶ Target optimized codegen
- ▶ Since IR depends on target, targets have to share the same or very similar ISA.

```
.section __TEXT,__text,regular,pure_instructions
.macosx_version_min 10, 11
.globl _main
.align 4, 0x90
_main:                                ## @main
    .cfi_startproc
## BB#0:
    push rbp
.Ltmp0:
    .cfi_def_cfa_offset 16
.Ltmp1:
    .cfi_offset rbp, -16
    mov rbp, rsp
.Ltmp2:
    .cfi_def_cfa_register rbp
    sub rsp, 16
    lea rdi, [rip + L_.str]
    mov al, 0
    call _printf
    xor ecx, ecx
    mov dword ptr [rbp - 4], eax ## 4-byte Spill
    mov eax, ecx
    add rsp, 16
    pop rbp
    ret
    .cfi_endproc

.section __TEXT,__cstring,cstring_literals
L_.str:                                ## @.str
    .asciz "Hello World"

.subsections_via_symbols
```

Bitcode Recompilation

- ▶ Bitcode as a format is nothing new. It's a bitstream representation of the LLVM-IR
- ▶ Goal: Provide micro-architecture optimizations for new devices from day 1.
 - ▶ Eg. NEON floating-point or hardware integer division
- ▶ Apple asks developers to upload Bitcode along with their binary, recompiles for new (micro-)architectures assuming they have the same ISA
- ▶ Separate from App Thinning

Bitcode Binaries (as uploaded to iTC)

- ▶ Fat archives with multiple architectures
- ▶ For each architecture, contains a specific Mach-O Segment __LLVM containing a section __bundle.
- ▶ The data segment contains a xar archive containing an LLVM Bitcode file per class

Extracting LLVM IR from Bitcode binaries

- ▶ [bitcode_retreiver](#) scans each architecture of the binary for the __LLVM segment and dumps it.
- ▶ Un-xar the Bitcode of the interested architecture (eg. xar -x -f arm.xar)
- ▶ Convert Bitcode Bitstream to human-readable LLVM-IR
- ▶ The LLVM IR can be obtained from the Bitcode by deserializing the bitstream with `llvm-dis`

Symbolication

- ▶ Application submitted with
 - ▶ .dSYM for compiled targets
 - ▶ .bcsymbolmap in case the Bitcode is recompiled, Apple will then provide new .dSYM files if recompiled
- ▶ Originally broken on launch according to some developers, now seems to be fixed.



Landon Fuller
@landonfuller

+ Follow

Bitcode breaking 3rd-party crash reporting already: binary UUID changes, and Apple doesn't offer dSYM for download.

iTunes Connect Opt-Out

- ▶ Opt-out possible on iOS
- ▶ Mandatory for tvOS and WatchOS apps
- ▶ ⚠ Apple might turn optimizations on or off with no control of developer
- ▶ ⚠ No information is provided about LLVM version or backend compilation flags

The screenshot shows the iTunes Connect interface for the app "Signal - Private Messenger". The top navigation bar includes "iTunes Connect My Apps", the app's logo, the developer name "Frederic Jacobs RIDDLE QUIET VENTURES, LLC", and links for "App Analytics" and "Sales and Trends". The main content area has tabs for "App Store", "Features", "TestFlight", and "Activity". On the left, under "APP STORE INFORMATION", there are sections for "App Information" and "Pricing and Availability", with the latter being the active tab. It lists two iOS app versions: "2.2 Waiting For Review" and "2.1.3 Ready for Sale". Below this is a "VERSION OR PLATFORM" section with a plus sign. The right side of the screen is titled "Pricing and Availability" and contains a section titled "Bitcode Auto-Recompilation". A note states: "Occasionally, we may automatically recompile apps that include bitcode to improve hardware support or to optimize our software. The checkbox below allows you to disable auto-recompilation for your app." A question "What happens if you disable bitcode auto-recompilation?" is followed by a bulleted list: "• Your app or a thinned version of your app may be unavailable for some devices.", "• Your app, and any app bundle that includes this app, may become unavailable whenever apps must be recompiled.", and "• If your app is unavailable on the App Store, Universal Purchase, redownloads, and Family Sharing won't work unless all platform versions have been approved.". Another section asks "To maintain your app's availability on the App Store:" with a bulleted list: "• Upload a new build of your app that contains bitcode.", "• Test the new build with TestFlight (optional).", and "• Submit that build with a new app version to App Review.". A link "Learn more about bitcode." is present, along with a checked checkbox "Don't use bitcode auto-recompilation." and a link "If you have any questions, contact us.". At the bottom, a link "Last-Compatible Version Settings" is visible.

SECURITY IMPLICATIONS



COMPILER OPTIMIZATIONS CORRECTNESS

Persistent State

- ▶ Data remains in memory beyond developer enforced boundaries
- ▶ Examples
 - ▶ Dead Store Elimination: store instruction is removed by the compiler because the result won't be read in subsequent instructions.
 - ▶ Function Call Inlining: Merges the stack frames of the caller and the callee. This results in changing the scope of variables stored on the function stack.
 - ▶ Code Motion: Reorders instructions and basic blocks in program based on their dependencies. Allows the compiler to reorder instructions and basic blocks in the program based on their dependencies

Dead Store Elimination

```
#include <string>
using std::string;

#include <memory>

// The specifics of this function are
// not important for demonstrating this bug.
const string getPasswordFromUser() const;

bool isPasswordCorrect() {
    bool isPasswordCorrect = false;
    string Password("password");

    if(Password == getPasswordFromUser()) {
        isPasswordCorrect = true;
    }

    // This line is removed from the optimized code
    // even though it secures the code by wiping
    // the password from memory.
    memset(Password, 0, sizeof(Password));

    return isPasswordCorrect;
}
```

https://gcc.gnu.org/bugzilla/show_bug.cgi?id=8537

Function Inlining

```
string decrypt(cipherText) {  
    // code fetches decryption key  
    // and decrypts message in secure environment  
}
```

```
// Fetches a message from the network  
// Deserialize the packet  
// Decrypt & parse result
```

```
void nextMessage() {  
    // Code in this function does not assume  
    // a trusted execution environment.  
    ...  
    //call secure function  
    protoBuf = decrypt(msg);  
    ...  
}
```

Undefined Behavior

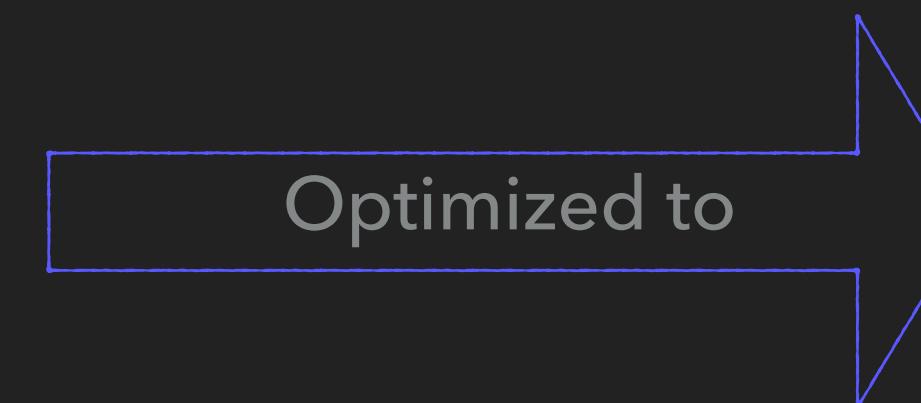
- ▶ Programming languages specifications intentionally leaving out semantic specifics for some operation
- ▶ Gives more freedom to implementors for optimizations and adapt to hardware
- ▶ GCC/Clang sometimes generate warnings for undefined behavior
- ▶ Examples in C:
 - ▶ uninitialized variable
 - ▶ dividing by zero
 - ▶ ...

Side Channels

- ▶ Any attack based on information gained from the physical implementation of a cryptosystem, rather than brute force or theoretical weaknesses in the algorithms
- ▶ Time and energy consumption are common and exploitable side-channels
- ▶ To prevent an attacker to guess what code path was taken, never branch on secret material

Example: Common Subexpression Elimination

```
int crypt(int k*){
    int key = 0;
    if (k[0]==0xC0DE){
        key=k[0]*15+3;
        key+=k[1]*15+3;
        key+=k[2]*15+3;
    } else {
        key=2*15+3;
        key+=2*15+3;
        key+=2*15+3;
    }
}
```



```
int crypt(int k*){
    int key = 0;
    if (k[0]==0xC0DE){
        key=k[0]*15+3;
        key+=k[1]*15+3;
        key+=k[2]*15+3;
    } else {
        tmp = 2*15+3;
        key = 3*tmp;
    }
}
```

⚠ Timing attack: One code path has less instructions than the other

**NO INLINE
ASSEMBLY**

There's just no way you can reasonably expect even the most advanced C compilers to do this on your behalf.

Mike Pall (LuaJIT author)
on possible optimizations

Inline Assembly

- ▶ Why?
- ▶ Gives more tuning power to developer to prevent side channel leakage
- ▶ Performance
- ▶ Widely used in cryptographic implementations (OpenSSL, Apple's own corecrypto ...)
- ▶ But not only, VLC also uses inline assembly for optimizations

COMPILER TRUST

How Much Do You Trust Your Compiler?

- ▶ Compromised compiler/toolchain
 - ▶ XcodeGhost, distribution of backdoor compiler
- ▶ Control over dependencies
 - ▶ A single “CocoaPod” module can compromise entire app
 - ▶ Relying on bad library
- ▶ Compiler bugs
- ▶ Buggy compiler/environment
- ▶ With Bitcode specifically, incentives for an attacker to compromise re-compilation is huge.
Exploit in build system could result in altering a huge amount of apps in App Store.

POTENTIAL FOR
BACKDOOR INJECTION

Obama Administration Working Group discussing using automatic updates to insert backdoors

Provider-enabled remote access to encrypted devices through current update procedures. Virtually all consumer devices include the capability to remotely download and install updates to their operating system and applications. For this approach, law enforcement would use lawful process to compel providers to use their remote update capability to insert law enforcement software into a targeted device. Once inserted, such software could enable far-reaching access to and control of the targeted device. This proposal would not require physical modification of devices, and so would likely be less costly for providers to implement. It would also enable remote access, and make surreptitious access much less costly. However, its use could call into question the trustworthiness of established software update channels. Individual users, concerned about remote access to their devices, could choose to turn off software updates, rendering their devices significantly less secure as time passed and vulnerabilities were discovered by not patched.

REPRODUCIBLE BUILDS

Reproducible Builds Without Bitcode

- ▶ Verifying that a build on the App Store matches the code on GitHub is hard
- ▶ Mostly due to FairPlay, Apple's FairPlay DRM
 - ▶ Can't opt out
 - ▶ Requires jailbroken iPhone, use lldb to dump binary at runtime

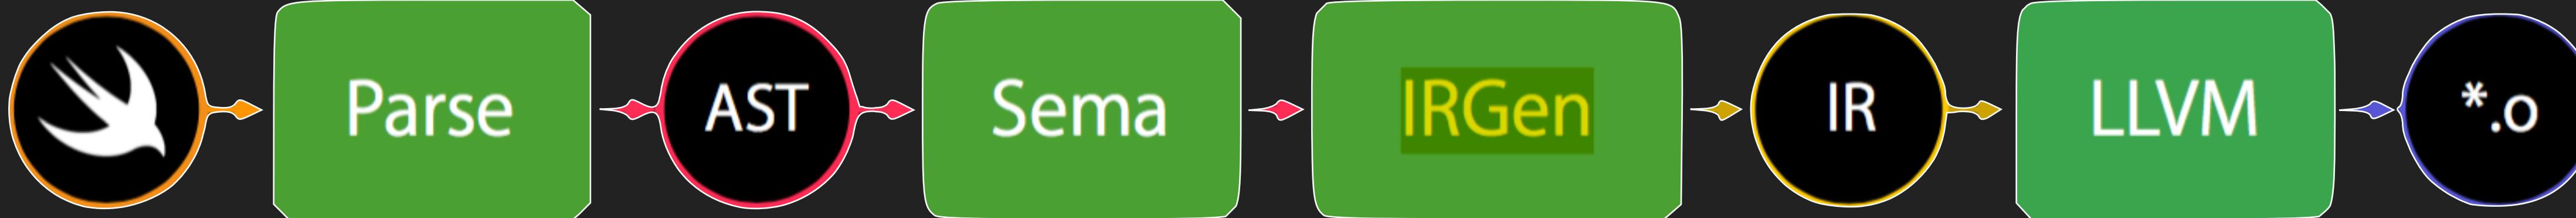
Reproducible Builds With Bitcode

- ▶ Same procedure for known architectures and if binaries were not recompiled
- ▶ But what happens when binaries don't match? Cost of reverse engineering differences is important if you don't have the toolchain

TOWARDS TARGET INDEPENDENCE?

Swift Intermediate Language

- ▶ Presented early November 2015 at the LLVM Dev Summit
- ▶ Higher-level IR
- ▶ Reducing undefined behavior
- ▶ IRGen handles type lowering & concretization of the target



QUESTIONS?

REFERENCES

- ▶ [Bitcode Demystified](#) by Alex Denisov
- ▶ Excellent resource on the [Correctness-Security Gap](#) from LangSec by Vijay D'Silva, Mathias Payer, Dawn Song
- ▶ Undefined behavior:
 - ▶ [Undefined Behavior: What Happened to My Code?](#)
 - ▶ <http://blog.llvm.org/2011/05/what-every-c-programmer-should-know.html>
 - ▶ <http://blog.regehr.org/archives/213>
- ▶ See clickable references in slides